

Name: Rio Marie G. Suzuki	Date Performed: 10/23/2023
Course/Section: CPE232S6	Date Submitted: 10/23/2023
Instructor: Dr. Jonathan Taylar	Semester and SY: 1st sem 2023-2024
Activity 10: Install, Configure, and Manage Log Monitoring tools	
1. Objectives	
Create and design a workflow that installs, configure and manage enterprise log monitoring tools using Ansible as an Infrastructure as Code (IaC) tool.	
2. Discussion	
<p>Log monitoring software scans and monitors log files generated by servers, applications, and networks. By detecting and alerting users to patterns in these log files, log monitoring software helps solve performance and security issues. System administrators use log monitoring software to detect common important events indicated by log files.</p> <p>Log monitoring software helps maintain IT infrastructure performance and pinpoints issues to prevent downtime and mitigate risks. These tools will often integrate with IT alerting software, log analysis software, and other IT issue resolution products to more aptly flesh out the IT infrastructure maintenance ecosystem.</p> <p>To qualify for inclusion in the Log Monitoring category, a product must:</p> <ul style="list-style-type: none"> • Monitor the log files generated by servers, applications, or networks • Alert users when important events are detected • Provide reporting capabilities for log files <p>Elastic Stack</p> <p>ELK suite stands for Elasticsearch, Kibana, Beats, and Logstash (also known as the ELK Stack). Source: https://www.elastic.co/elastic-stack</p> <p>The Elastic Stack is a group of open source products from Elastic designed to help users take data from any type of source and in any format, and search, analyze and visualize that data in real time. The product group was formerly known as the ELK Stack for the core products in the group -- Elasticsearch, Logstash and Kibana -- but has been rebranded as the Elastic Stack. A fourth product, Beats, was subsequently added to the stack. The Elastic Stack can be deployed on premises or made available as software as a service (SaaS). Elasticsearch supports Amazon Web Services (AWS), Google Cloud Platform and Microsoft Azure.</p> <p>GrayLog</p> <p>Graylog is a powerful platform that allows for easy log management of both structured and unstructured data along with debugging applications.</p>	

It is based on Elasticsearch, MongoDB, and Scala. Graylog has a main server, which receives data from its clients installed on different servers, and a web interface, which visualizes the data and allows to work with logs aggregated by the main server.

We use Graylog primarily as the stash for the logs of the web applications we build. However, it is also effective when working with raw strings (i.e. syslog): the tool parses it into the structured data we need. It also allows advanced custom search in the logs using structured queries. In other words, when integrated properly with a web app, Graylog helps engineers to analyze the system behavior on almost per code line basis.

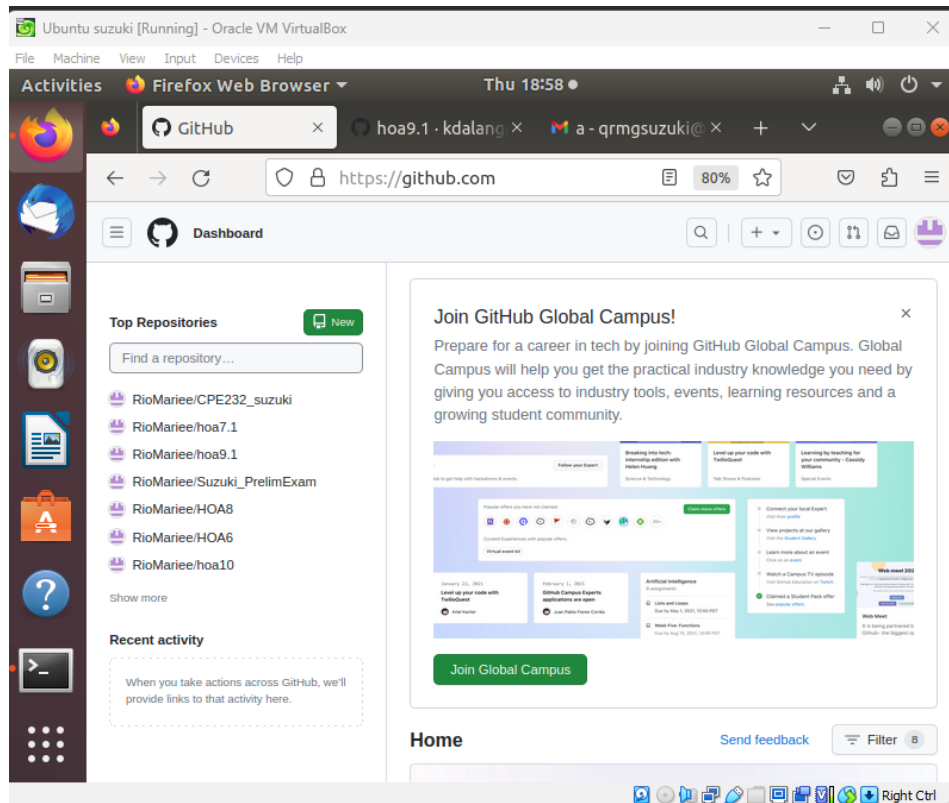
Source: <https://www.graylog.org/products/open-source>

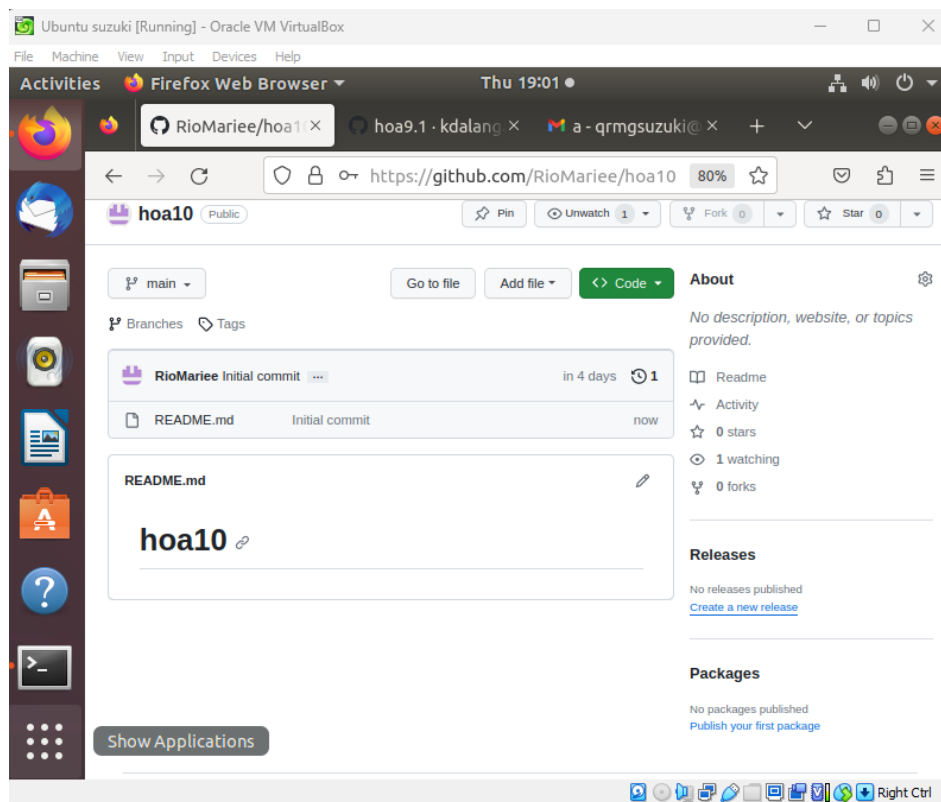
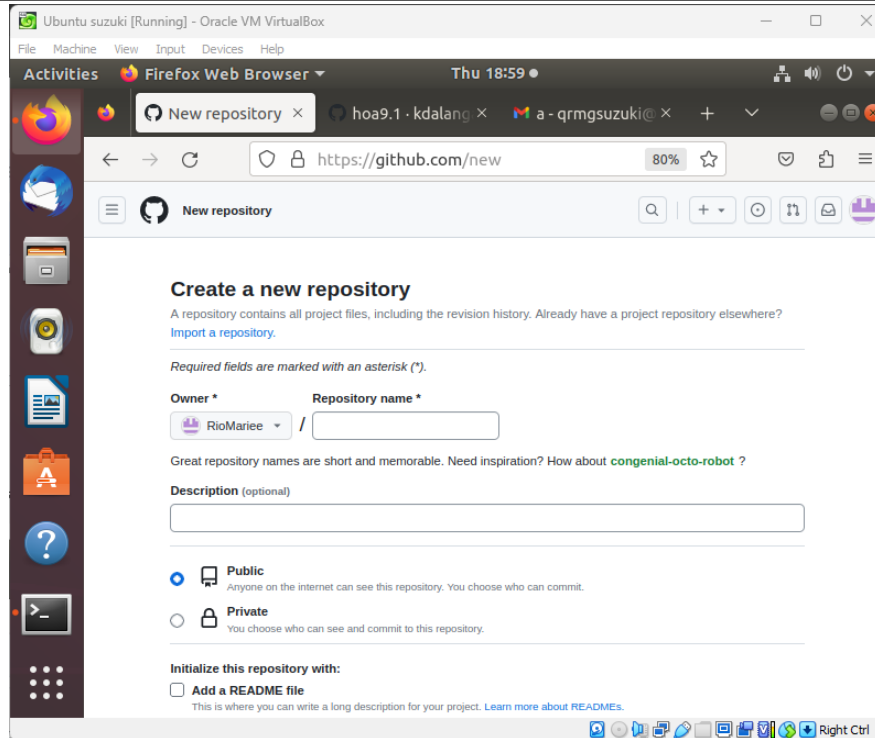
3. Tasks

1. Create a playbook that:
 - a. Install and configure Elastic Stack in separate hosts (Elastic Search, Kibana, Logstash)
2. Apply the concept of creating roles.
3. Describe how you did step 1. (Provide screenshots and explanations in your report. Make your report detailed such that it will look like a manual.)
4. Show an output of the installed Elastic Stack for both Ubuntu and CentOS.
5. Make sure to create a new repository in GitHub for this activity.

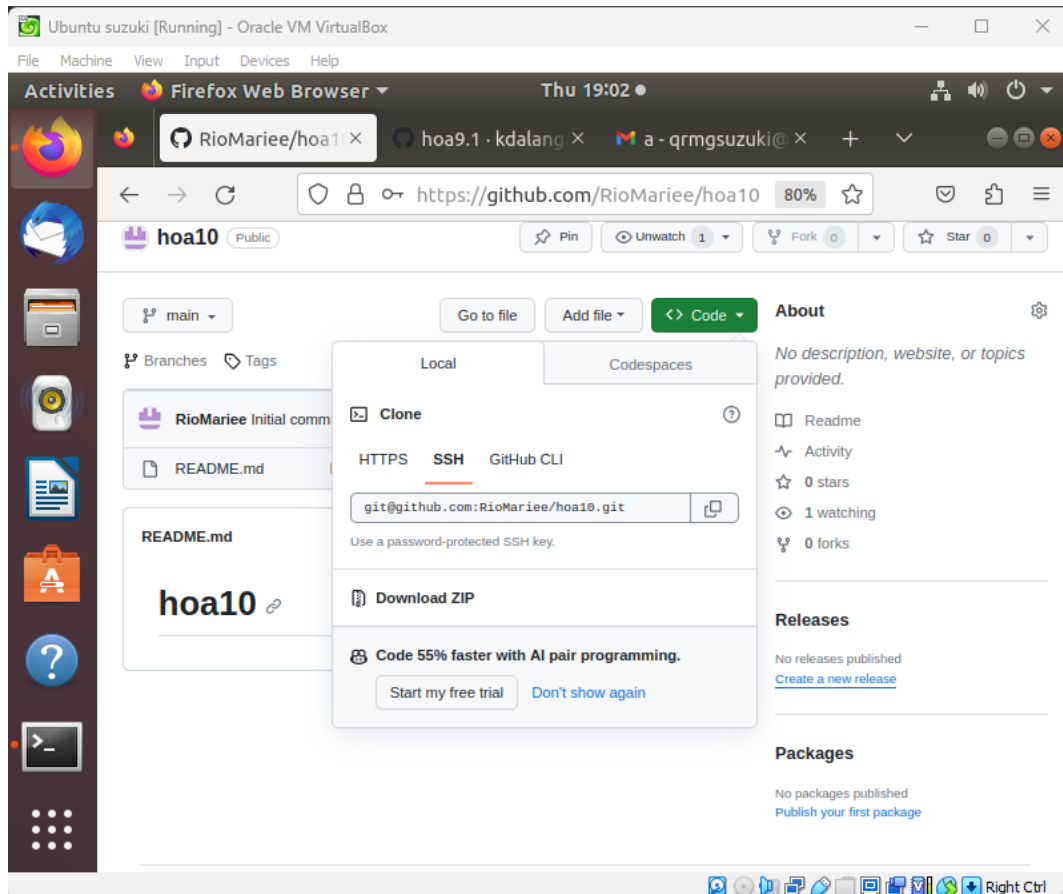
4. Output (screenshots and explanations)

Step 1. First thing first, we will create a new repository for this activity. To create a new repository, we will go to git hub.com and sign in our account. On the left side you will see the list of repositories and click create new repository and name it “hoa10”.





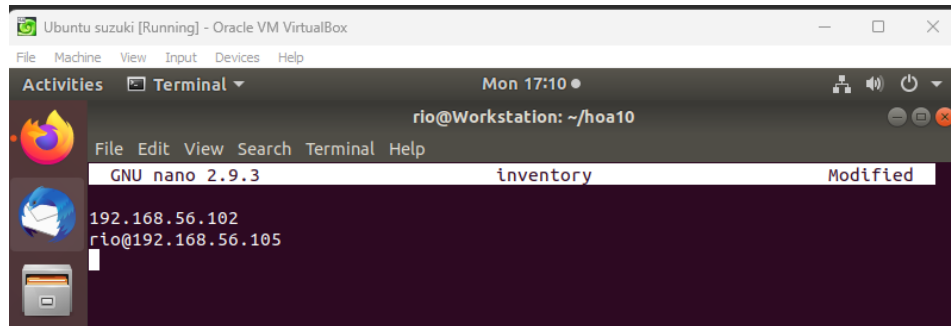
Step 2. Next step we will copy the link of the created new repository by clicking the code and select “ssh” then click the copy button. After copying the ssh we will now paste it in the Ubuntu terminal together with the command “git clone (ssh of the repository)”.



Output:

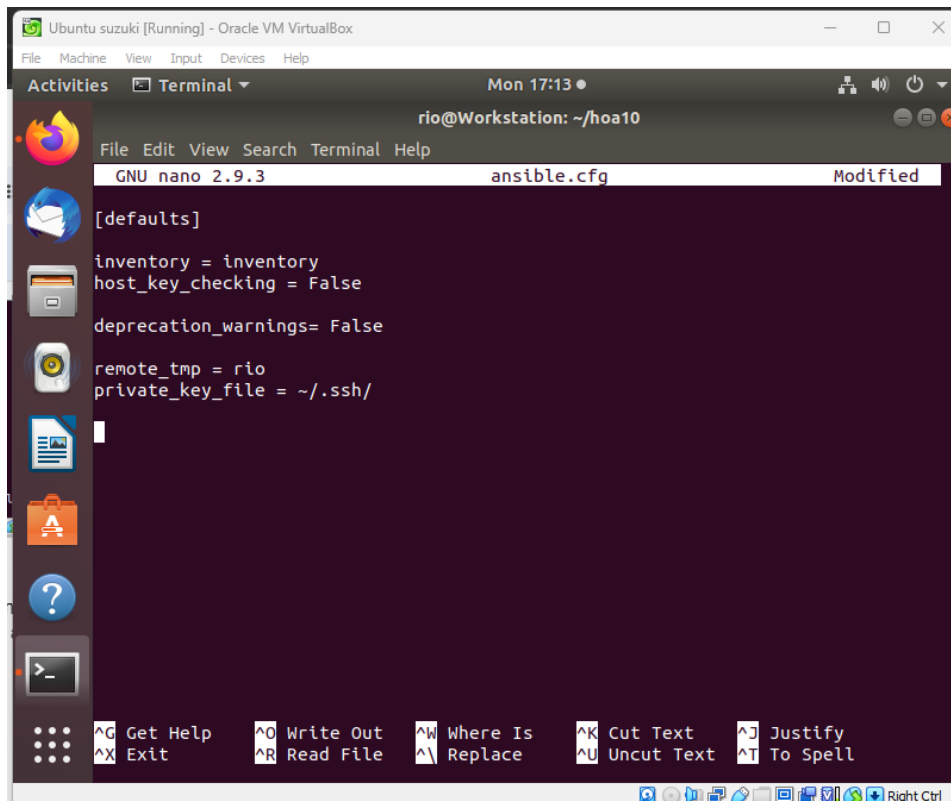
```
File Edit View Search Terminal Help
rio@Workstation: ~
rio@Workstation:~$ git clone git@github.com:RioMarieee/hoa10.git
Cloning into 'hoa10'...
remote: Enumerating objects: 3, done.
remote: Counting objects: 100% (3/3), done.
remote: Total 3 (delta 0), reused 0 (delta 0), pack-reused 0
Receiving objects: 100% (3/3), done.
rio@Workstation:~$
```

Step 3. After cloning the repository we will now enter the repository we created in Ubuntu and create the inventory using the command “sudo nano inventory” and put the ip address of the centos and the server 1 ip address, and don’t forget to save it.



The screenshot shows a terminal window titled "Ubuntu suzuki [Running] - Oracle VM VirtualBox". The terminal prompt is "rio@Workstation: ~/hoa10". A nano editor window is open, editing a file named "inventory". The file contains two lines of IP addresses: "192.168.56.102" and "rio@192.168.56.105". The nano editor status bar at the top indicates "GNU nano 2.9.3" and "Modified".

Step 4. After creating the inventory we will now also create the ansible.cfg. The contents of the ansible.cfg will remain the same just like in the previous activities. Press ctrl+z then press y to confirm the save settings.

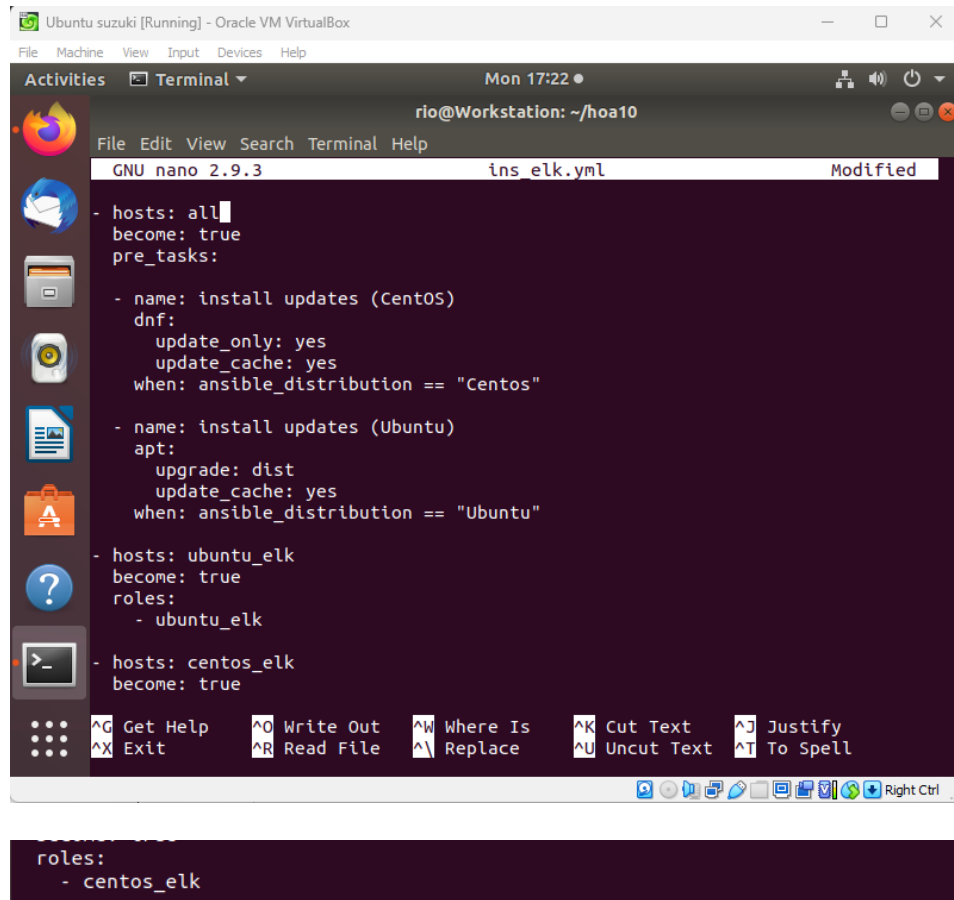


The screenshot shows a terminal window titled "Ubuntu suzuki [Running] - Oracle VM VirtualBox". The terminal prompt is "rio@Workstation: ~/hoa10". A nano editor window is open, editing a file named "ansible.cfg". The file contains the following configuration:

```
[defaults]
inventory = inventory
host_key_checking = False
deprecation_warnings= False
remote_tmp = rio
private_key_file = ~/.ssh/
```

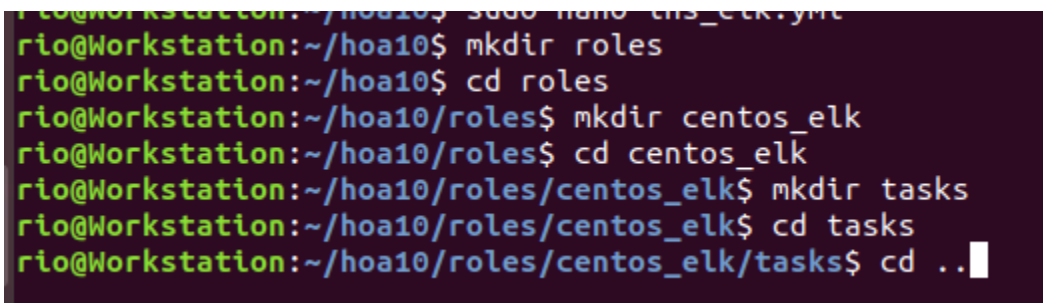
The nano editor status bar at the top indicates "GNU nano 2.9.3" and "Modified". At the bottom of the terminal window, there is a status bar with various keyboard shortcuts: ^G Get Help, ^O Write Out, ^W Where Is, ^K Cut Text, ^J Justify, ^X Exit, ^R Read File, ^_ Replace, ^U Uncut Text, ^T To Spell.

Step 5. Next we will create the `ins_elk.yml`, this will contain the updates needed for both Ubuntu and CentOS. After creating the `ins_elk.yml` don't forget to save it.



```
GNU nano 2.9.3 ins_elk.yml Modified
- hosts: all
  become: true
  pre_tasks:
    - name: install updates (CentOS)
      dnf:
        update_only: yes
        update_cache: yes
      when: ansible_distribution == "Centos"
    - name: install updates (Ubuntu)
      apt:
        upgrade: dist
        update_cache: yes
      when: ansible_distribution == "Ubuntu"
- hosts: ubuntu_elk
  become: true
  roles:
    - ubuntu_elk
- hosts: centos_elk
  become: true
  roles:
    - centos_elk
```

Step 6. After we created the `ins_elk.yml`, we will now create a tree consisting of roles and `centos_elk` and `ubuntu_elk` under each of them will be the tasks and the `main.yml`. to create the tree we will use the command “`mkdir (name of the directory)`” and the command “`cd (name of the created new directory)`” to change directory, and “`cd ..`” to go back.

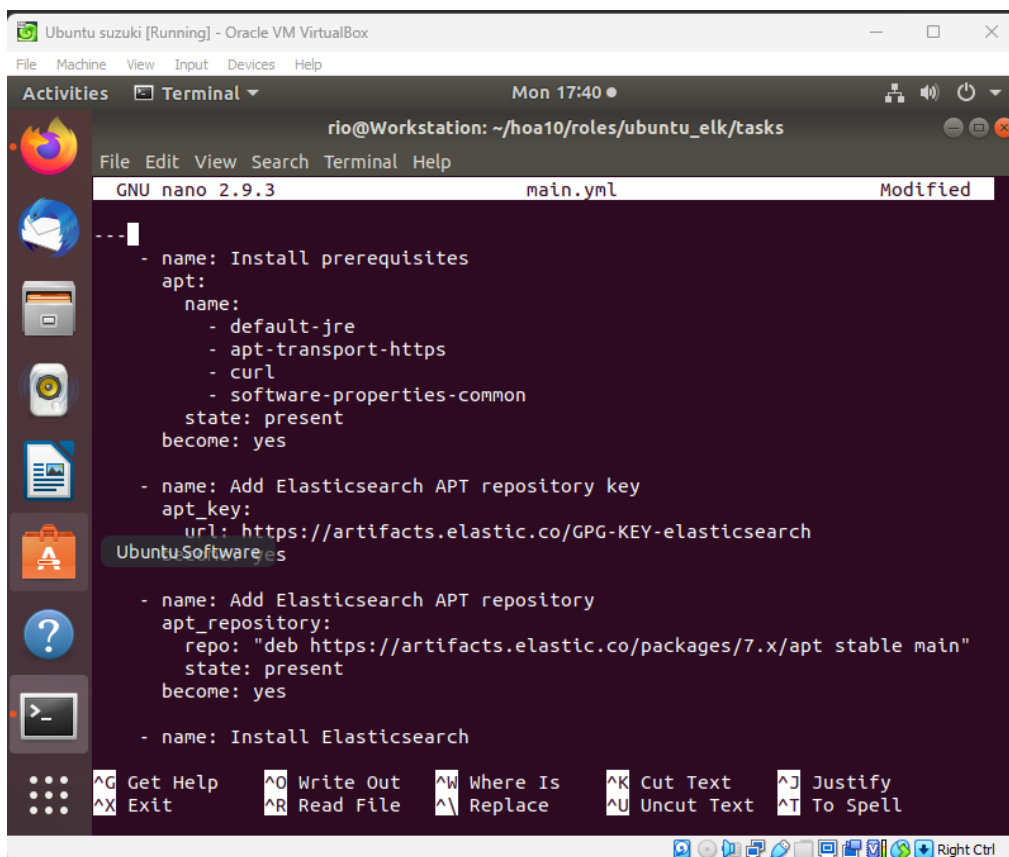


```
rio@Workstation:~/hoa10$ mkdir roles
rio@Workstation:~/hoa10$ cd roles
rio@Workstation:~/hoa10/roles$ mkdir centos_elk
rio@Workstation:~/hoa10/roles$ cd centos_elk
rio@Workstation:~/hoa10/roles/centos_elk$ mkdir tasks
rio@Workstation:~/hoa10/roles/centos_elk$ cd tasks
rio@Workstation:~/hoa10/roles/centos_elk/tasks$ cd ..
```

```
rio@Workstation:~/hoa10/roles$ cd ubuntu_elk
rio@Workstation:~/hoa10/roles/ubuntu_elk$ mkdir tasks
rio@Workstation:~/hoa10/roles/ubuntu_elk$ cd tasks
rio@Workstation:~/hoa10/roles/ubuntu_elk/tasks$ sudo nano main.yml
```

Step 7. After creating the `ubuntu_elk` and the `centos_elk`, we will now create the `main.yml` under the `ubuntu_elk` and `centos_elk`. Also, don't forget to save your work.

Input: `Ubuntu_elk` `main.yml` contents:



The screenshot shows a terminal window titled "Ubuntu suzuki [Running] - Oracle VM VirtualBox". The terminal is running the `nano` text editor to edit `main.yml`. The file contains the following YAML content:

```
---
- name: Install prerequisites
  apt:
    name:
      - default-jre
      - apt-transport-https
      - curl
      - software-properties-common
    state: present
    become: yes

- name: Add Elasticsearch APT repository key
  apt_key:
    url: https://artifacts.elastic.co/GPG-KEY-elasticsearch
    UbuntuSoftwarees

- name: Add Elasticsearch APT repository
  apt_repository:
    repo: "deb https://artifacts.elastic.co/packages/7.x/apt stable main"
    state: present
    become: yes

- name: Install Elasticsearch
```

The terminal window also shows a sidebar with application icons and a bottom status bar with system icons and a "Right Ctrl" indicator.

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 17:41

rio@Workstation: ~/hoa10/roles/ubuntu_elk/tasks

File Edit View Search Terminal Help

GNU nano 2.9.3 main.yml Modified

```
apt_key:
  url: https://artifacts.elastic.co/GPG-KEY-elasticsearch
  become: yes

- name: Add Elasticsearch APT repository
  apt_repository:
    repo: "deb https://artifacts.elastic.co/packages/7.x/apt stable main"
    state: present
    become: yes

- name: Install Elasticsearch
  apt:
    name: elasticsearch
    state: present
    become: yes

- name: Enable and start Elasticsearch service
  systemd:
    name: elasticsearch
    enabled: yes
    state: started
    become: yes
```

Get Help Write Out Where Is Cut Text Justify
Exit Read File Replace Uncut Text To Spell

Right Ctrl

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 17:41

rio@Workstation: ~/hoa10/roles/ubuntu_elk/tasks

File Edit View Search Terminal Help

GNU nano 2.9.3 main.yml Modified

```
name: elasticsearch
state: present
become: yes

- name: Enable and start Elasticsearch service
  systemd:
    name: elasticsearch
    enabled: yes
    state: started
    become: yes

- name: Install Kibana
  apt:
    name: kibana
    state: present
    become: yes

- name: Enable and start Kibana service
  systemd:
    name: kibana
    enabled: yes
    state: started
    become: yes
```

Get Help Write Out Where Is Cut Text Justify
Exit Read File Replace Uncut Text To Spell

Right Ctrl

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 17:41

rio@Workstation: ~/hoa10/roles/ubuntu_elk/tasks

File Edit View Search Terminal Help

GNU nano 2.9.3 main.yml Modified

```
apt:
  name: kibana
  state: present
  become: yes

- name: Enable and start Kibana service
  systemd:
    name: kibana
    enabled: yes
    state: started
    become: yes

- name: Install Logstash
  apt:
    name: logstash
    state: present
    become: yes

- name: Enable and start Logstash service
  systemd:
    name: logstash
    enabled: yes
    state: started
```

Get Help Write Out Where Is Cut Text Justify
Exit Read File Replace Uncut Text To Spell

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 17:42

rio@Workstation: ~/hoa10/roles/ubuntu_elk/tasks

File Edit View Search Terminal Help

GNU nano 2.9.3 main.yml Modified

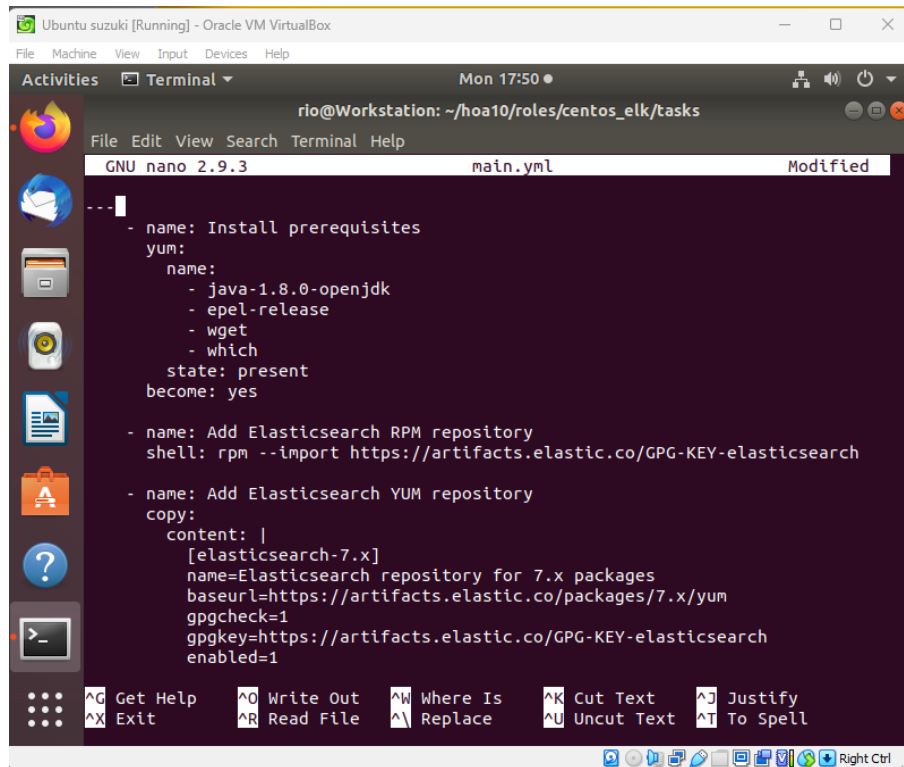
```
- name: Install Logstash
  apt:
    name: logstash
    state: present
    become: yes

- name: Enable and start Logstash service
  systemd:
    name: logstash
    enabled: yes
    state: started
    become: yes

- name: Restart Elasticsearch and Kibana
  systemd:
    name: "{{ item }}"
    state: restarted
  loop:
    - elasticsearch
    - kibana
```

Get Help Write Out Where Is Cut Text Justify
Exit Read File Replace Uncut Text To Spell

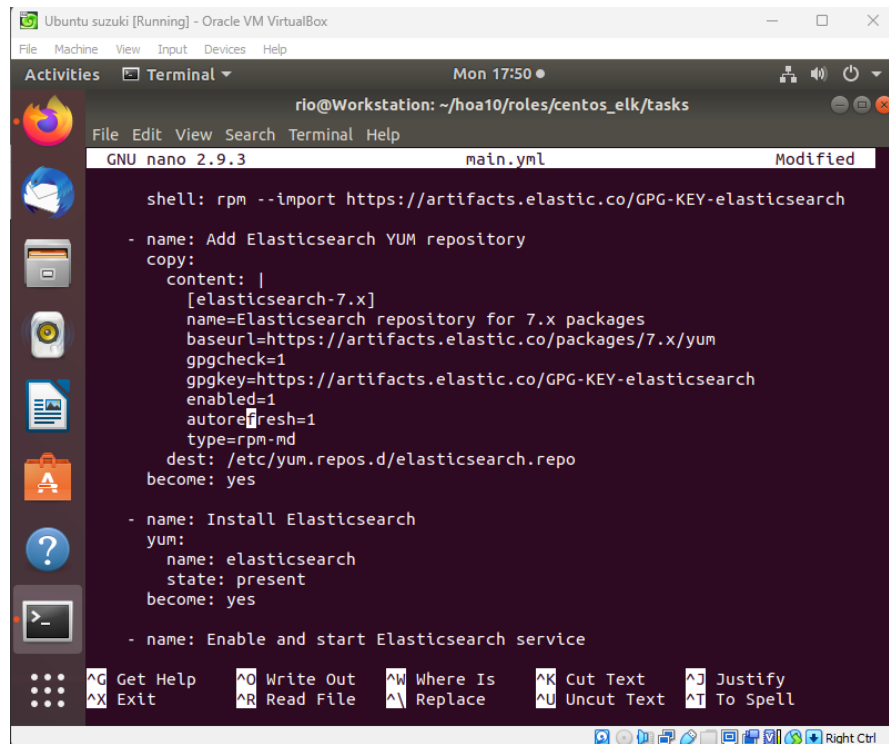
Input: contents of centos_elk main.yml.



The screenshot shows a terminal window titled 'Ubuntu suzuki [Running] - Oracle VM VirtualBox'. The terminal is running the 'nano' text editor, editing a file named 'main.yml'. The user is 'rio@Workstation' and the current directory is '~/hoa10/roles/centos_elk/tasks'. The file content is as follows:

```
---  
- name: Install prerequisites  
  yum:  
    name:  
      - java-1.8.0-openjdk  
      - epel-release  
      - wget  
      - which  
    state: present  
    become: yes  
  
- name: Add Elasticsearch RPM repository  
  shell: rpm --import https://artifacts.elastic.co/GPG-KEY-elasticsearch  
  
- name: Add Elasticsearch YUM repository  
  copy:  
    content: |  
      [elasticsearch-7.x]  
      name=Elasticsearch repository for 7.x packages  
      baseurl=https://artifacts.elastic.co/packages/7.x/yum  
      gpgcheck=1  
      gpgkey=https://artifacts.elastic.co/GPG-KEY-elasticsearch  
      enabled=1
```

The terminal window includes a menu bar (File, Edit, View, Search, Terminal, Help) and a status bar at the bottom with keyboard shortcuts for various nano editor functions.



This screenshot shows the continuation of the 'main.yml' file in the nano editor. The content continues from the previous state:

```
      shell: rpm --import https://artifacts.elastic.co/GPG-KEY-elasticsearch  
  
- name: Add Elasticsearch YUM repository  
  copy:  
    content: |  
      [elasticsearch-7.x]  
      name=Elasticsearch repository for 7.x packages  
      baseurl=https://artifacts.elastic.co/packages/7.x/yum  
      gpgcheck=1  
      gpgkey=https://artifacts.elastic.co/GPG-KEY-elasticsearch  
      enabled=1  
      autorefresh=1  
      type=rpm-md  
      dest: /etc/yum.repos.d/elasticsearch.repo  
      become: yes  
  
- name: Install Elasticsearch  
  yum:  
    name: elasticsearch  
    state: present  
    become: yes  
  
- name: Enable and start Elasticsearch service
```

The terminal window maintains the same interface as the previous screenshot, showing the user's progress in editing the configuration file.

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 17:50

rio@Workstation: ~/hoa10/roles/centos_elk/tasks

File Edit View Search Terminal Help

GNU nano 2.9.3 main.yml Modified

```
type=rpm-md
dest: /etc/yum.repos.d/elasticsearch.repo
become: yes

- name: Install Elasticsearch
  yum:
    name: elasticsearch
    state: present
    become: yes

- name: Enable and start Elasticsearch service
  systemd:
    name: elasticsearch
    enabled: yes
    state: started
    become: yes

- name: Install Kibana
  yum:
    name: kibana
    state: present
    become: yes
```

Get Help Write Out Where Is Cut Text Justify
Exit Read File Replace Uncut Text To Spell

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 17:51

rio@Workstation: ~/hoa10/roles/centos_elk/tasks

File Edit View Search Terminal Help

GNU nano 2.9.3 main.yml Modified

```
name: elasticsearch
enabled: yes
state: started
become: yes

- name: Install Kibana
  yum:
    name: kibana
    state: present
    become: yes

- name: Enable and start Kibana service
  systemd:
    name: kibana
    enabled: yes
    state: started
    become: yes

- name: Install Logstash
  yum:
    name: logstash
    state: present
    become: yes
```

Get Help Write Out Where Is Cut Text Justify
Exit Read File Replace Uncut Text To Spell

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 17:51

rio@Workstation: ~/hoa10/roles/centos_elk/tasks

File Edit View Search Terminal Help

GNU nano 2.9.3 main.yml Modified

```
systemd:
  name: kibana
  enabled: yes
  state: started
  become: yes

- name: Install Logstash
  yum:
    name: logstash
    state: present
    become: yes

- name: Enable and start Logstash service
  systemd:
    name: logstash
    enabled: yes
    state: started
    become: yes

- name: Restart Elasticsearch and Kibana
  systemd:
    name: "{{ item }}"
    state: restarted
```

Terminal

Get Help Write Out Where Is Cut Text Justify
Exit Read File Replace Uncut Text To Spell

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 17:51

rio@Workstation: ~/hoa10/roles/centos_elk/tasks

File Edit View Search Terminal Help

GNU nano 2.9.3 main.yml Modified

```
- name: Enable and start Logstash service
  systemd:
    name: logstash
    enabled: yes
    state: started
    become: yes

- name: Restart Elasticsearch and Kibana
  systemd:
    name: "{{ item }}"
    state: restarted
  loop:
    - elasticsearch
    - kibana
```

Terminal

Get Help Write Out Where Is Cut Text Justify
Exit Read File Replace Uncut Text To Spell

Step 8. After creating the main.yml of the activities we will now issue the command “tree” to double check if we did the tree right. To exit the directory you are currently in use the command “cd ..”.

```
rio@Workstation:~/hoa10/roles/centos_elk/tasks$ cd ..
rio@Workstation:~/hoa10/roles/centos_elk$ cd ..
rio@Workstation:~/hoa10/roles$ cd ..
rio@Workstation:~/hoa10$
```

```
rio@Workstation:~/hoa10$ tree
```

```

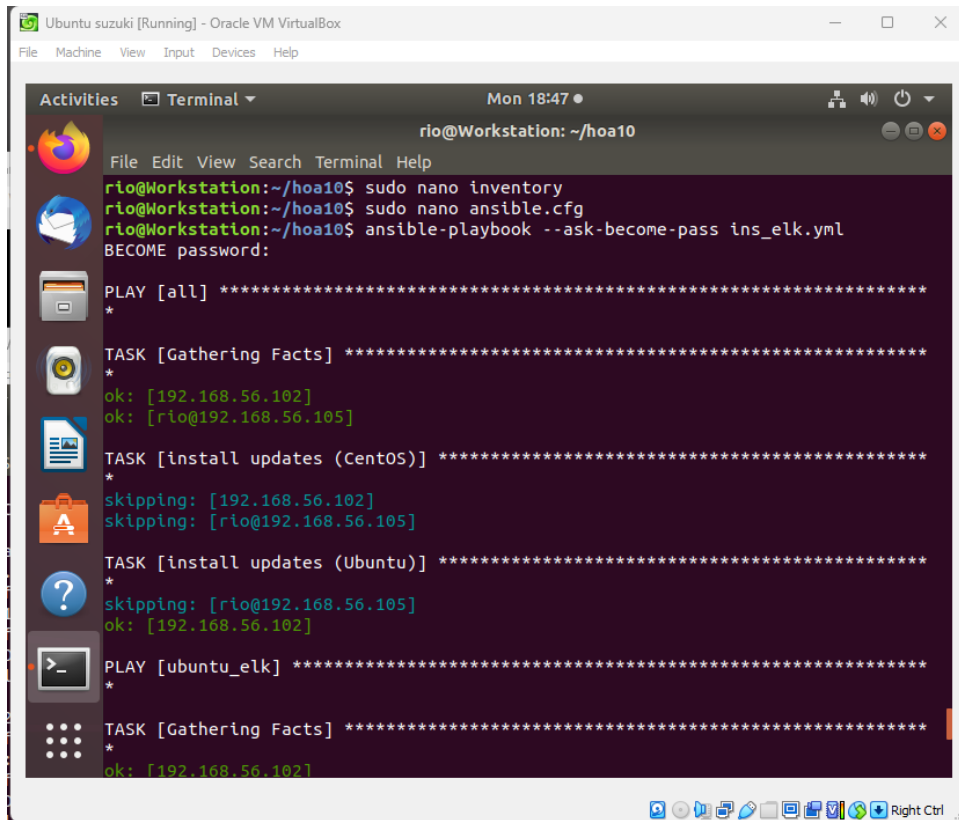
├── ansible.cfg
├── ins_elk.yml
├── inventory
├── README.md
└── roles
    ├── centos_elk
    │   └── tasks
    │       └── main.yml
    └── ubuntu_elk
        └── tasks
            └── main.yml
```

```
5 directories, 6 files
```

```
rio@Workstation:~/hoa10$
```

Step 9. Now run the command “Ansible-playbook –ask-become-pass ins_elk” and enter your workstation password this will run the commands we have in the ins_elk and if it has no error then the command is right and working.

Process:



Terminal window showing the execution of an Ansible playbook. The user is logged in as 'rio' on a workstation named 'hoa10'. The terminal displays the following commands and output:

```
rio@Workstation: ~/hoa10
File Edit View Search Terminal Help
rio@Workstation:~/hoa10$ sudo nano inventory
rio@Workstation:~/hoa10$ sudo nano ansible.cfg
rio@Workstation:~/hoa10$ ansible-playbook --ask-become-pass ins_elk.yml
BECOME password:

PLAY [all] *****
*

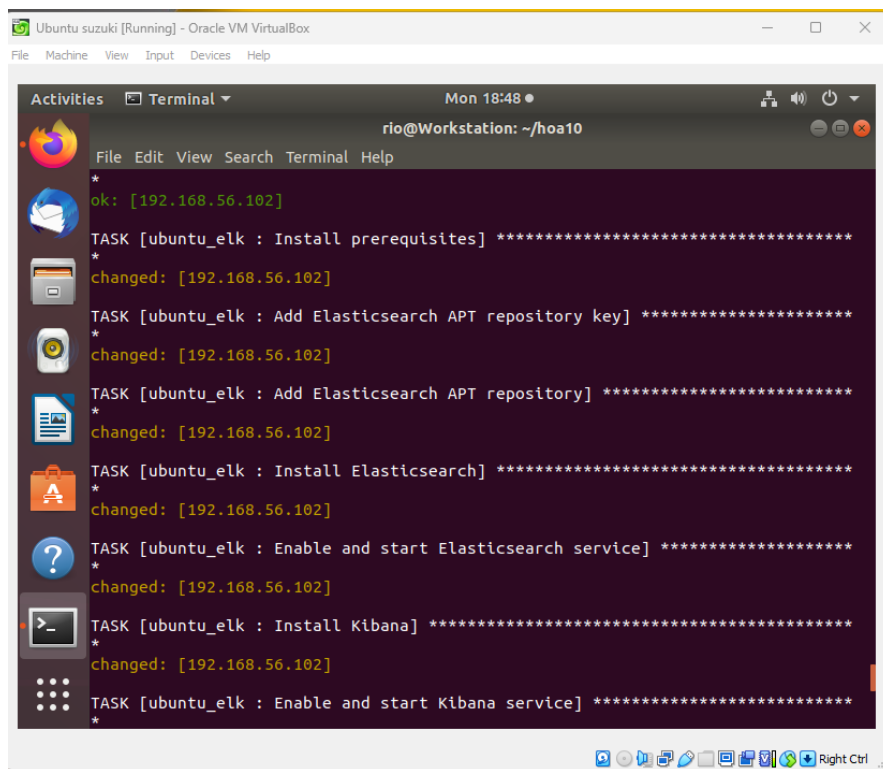
TASK [Gathering Facts] *****
*
ok: [192.168.56.102]
ok: [rio@192.168.56.105]

TASK [install updates (CentOS)] *****
*
skipping: [192.168.56.102]
skipping: [rio@192.168.56.105]

TASK [install updates (Ubuntu)] *****
*
skipping: [rio@192.168.56.105]
ok: [192.168.56.102]

PLAY [ubuntu_elk] *****
*

TASK [Gathering Facts] *****
*
ok: [192.168.56.102]
```



Terminal window showing the continuation of the Ansible playbook execution. The user is logged in as 'rio' on a workstation named 'hoa10'. The terminal displays the following commands and output:

```
rio@Workstation: ~/hoa10
File Edit View Search Terminal Help
*
ok: [192.168.56.102]

TASK [ubuntu_elk : Install prerequisites] *****
*
changed: [192.168.56.102]

TASK [ubuntu_elk : Add Elasticsearch APT repository key] *****
*
changed: [192.168.56.102]

TASK [ubuntu_elk : Add Elasticsearch APT repository] *****
*
changed: [192.168.56.102]

TASK [ubuntu_elk : Install Elasticsearch] *****
*
changed: [192.168.56.102]

TASK [ubuntu_elk : Enable and start Elasticsearch service] *****
*
changed: [192.168.56.102]

TASK [ubuntu_elk : Install Kibana] *****
*
changed: [192.168.56.102]

TASK [ubuntu_elk : Enable and start Kibana service] *****
*
```

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 18:49 rio@Workstation: ~/hoa10

```
File Edit View Search Terminal Help
changed: [192.168.56.102]

TASK [ubuntu_elk : Install Logstash] *****
*
changed: [192.168.56.102]

TASK [ubuntu_elk : Enable and start Logstash service] *****
*
changed: [192.168.56.102]

TASK [ubuntu_elk : Restart Elasticsearch and Kibana] *****
*
changed: [192.168.56.102] => (item=elasticsearch)
changed: [192.168.56.102] => (item=kibana)

PLAY [centos_elk] *****
*

TASK [Gathering Facts] *****
*
ok: [rio@192.168.56.105]

TASK [centos_elk : Install prerequisites] *****
*
ok: [rio@192.168.56.105]

TASK [centos_elk : Add Elasticsearch RPM repository] *****
*
changed: [rio@192.168.56.105]
```

Right Ctrl

Ubuntu suzuki [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Activities Terminal Mon 19:02 rio@Workstation: ~/hoa10

```
File Edit View Search Terminal Help

TASK [centos_elk : Install prerequisites] *****
*
ok: [rio@192.168.56.105]

TASK [centos_elk : Add Elasticsearch RPM repository] *****
*
changed: [rio@192.168.56.105]

TASK [centos_elk : Add Elasticsearch YUM repository] *****
*
changed: [rio@192.168.56.105]

TASK [centos_elk : Install Elasticsearch] *****
*
changed: [rio@192.168.56.105]

TASK [centos_elk : Enable and start Elasticsearch service] *****
*
changed: [rio@192.168.56.105]

TASK [centos_elk : Install Kibana] *****
*
changed: [rio@192.168.56.105]

TASK [centos_elk : Enable and start Kibana service] *****
*
changed: [rio@192.168.56.105]

TASK [centos_elk : Install Loostash] *****
```

Right Ctrl

The screenshot shows a terminal window titled 'Ubuntu suzuki [Running] - Oracle VM VirtualBox'. The terminal is running a playbook on a host named 'rio@Workstation' with IP '192.168.56.105'. The playbook tasks include enabling and starting Kibana, installing Logstash, enabling and starting Logstash, and restarting Elasticsearch and Kibana. The output shows that all tasks were successful, with no errors or warnings.

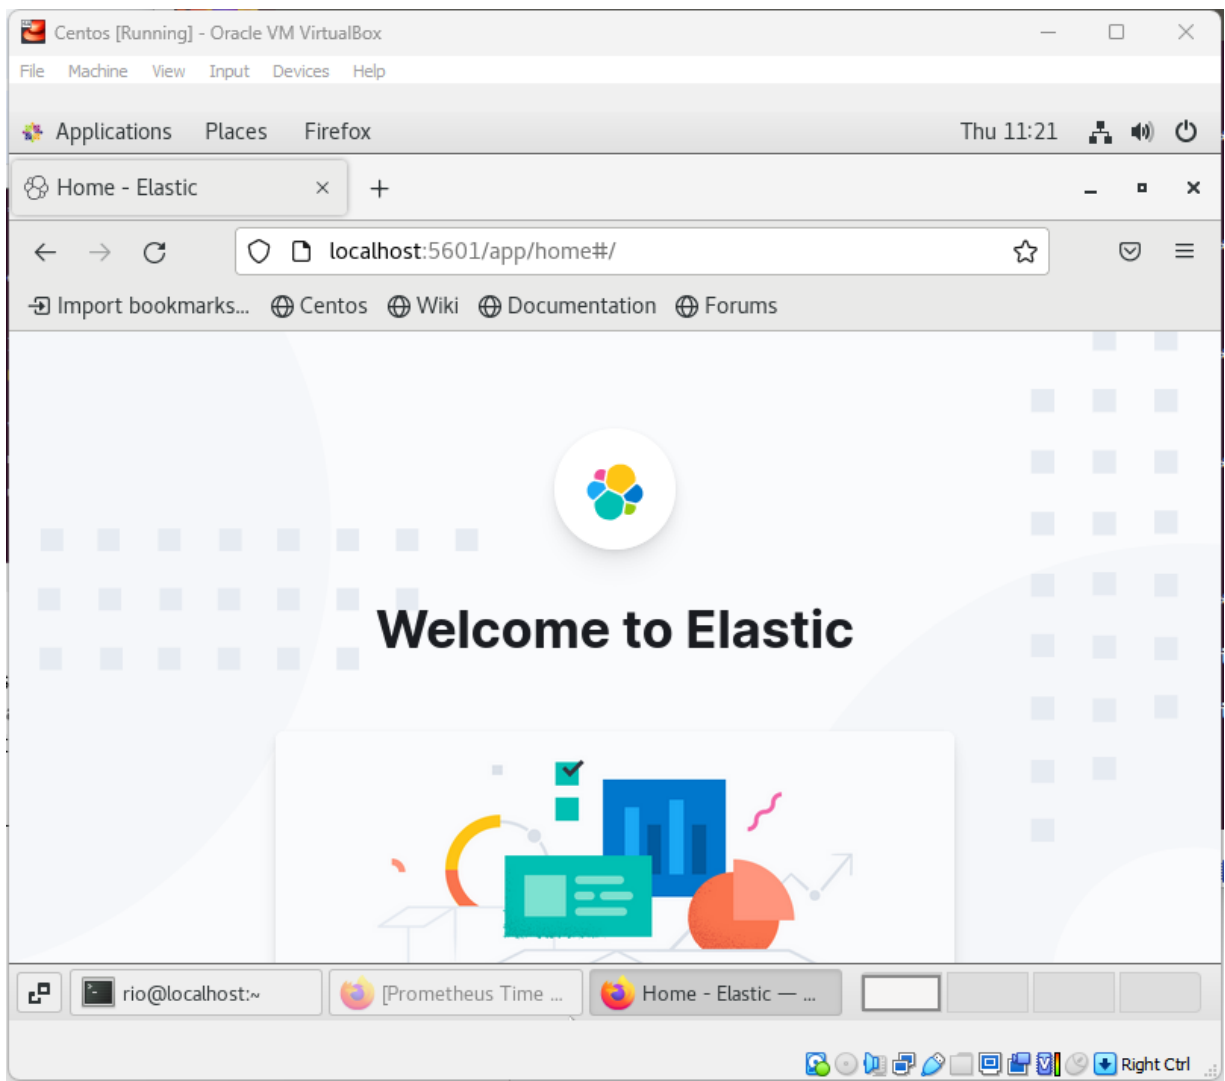

```
File Edit View Search Terminal Help
*
changed: [rio@192.168.56.105]
TASK [centos_elk : Enable and start Kibana service] *****
*
changed: [rio@192.168.56.105]
TASK [centos_elk : Install Logstash] *****
*
changed: [rio@192.168.56.105]
TASK [centos_elk : Enable and start Logstash service] *****
*
changed: [rio@192.168.56.105]
TASK [centos_elk : Restart Elasticsearch and Kibana] *****
*
changed: [rio@192.168.56.105] => (item=elasticsearch)
changed: [rio@192.168.56.105] => (item=kibana)
PLAY RECAP *****
192.168.56.102 : ok=13 changed=10 unreachable=0 failed=0
skipped=1 rescued=0 ignored=0
rio@192.168.56.105 : ok=12 changed=9 unreachable=0 failed=0
skipped=2 rescued=0 ignored=0
rio@Workstation:~/hoa10$
```

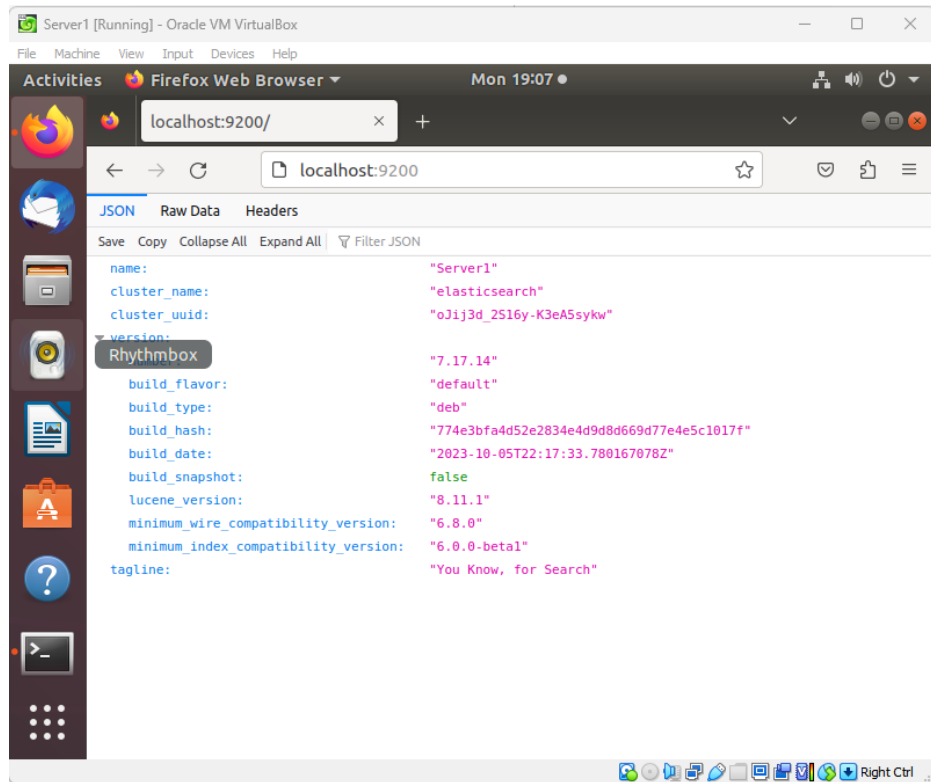
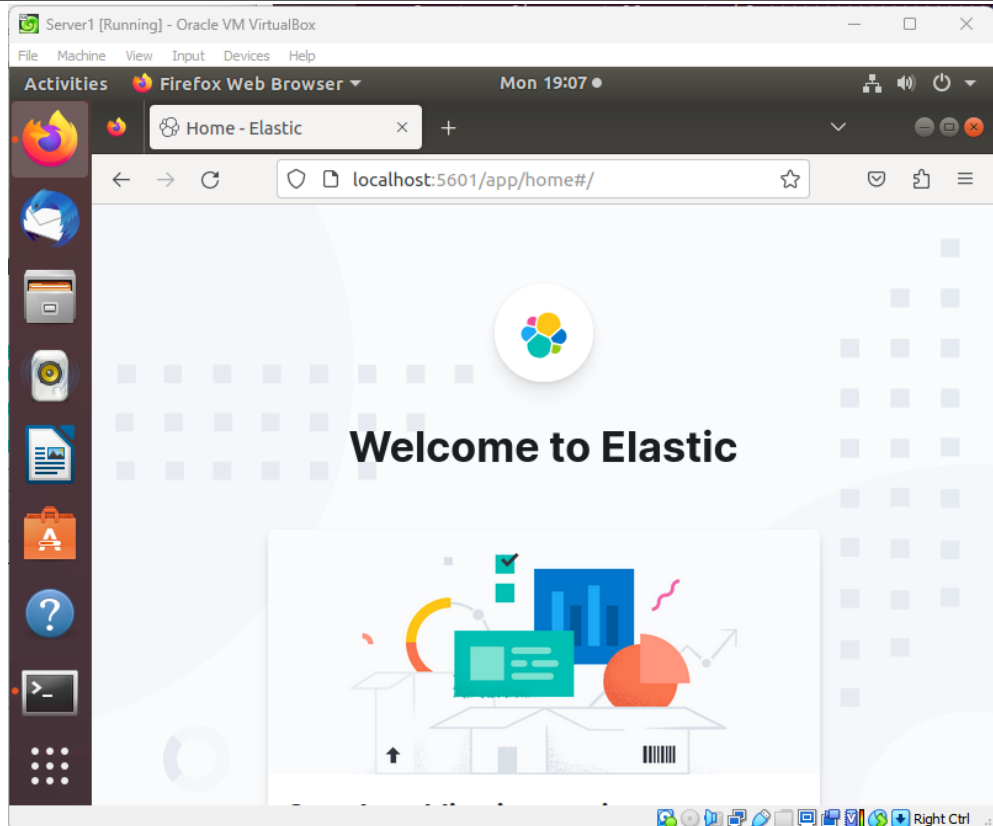
Step 10. We will now check the evidence of the successful installation using the firefox of the centos and ubuntu. To check if we have the successfully install the packages we will use the command “localhost:5601” it should redirect to

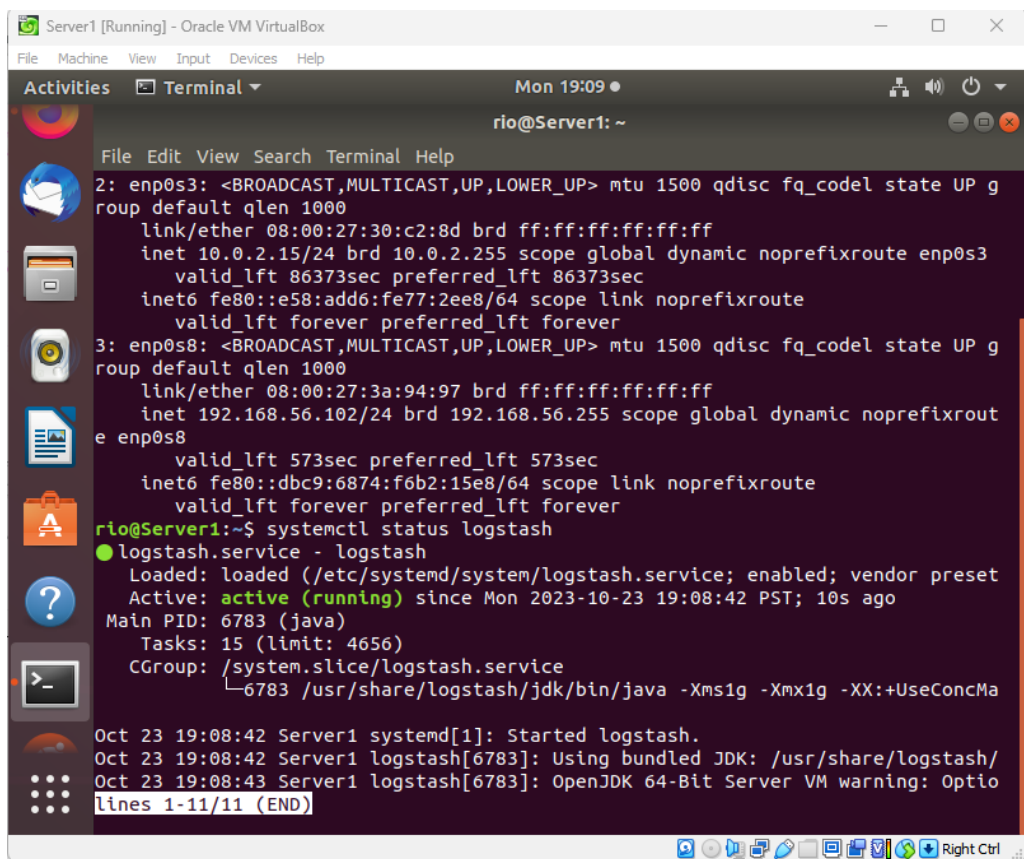
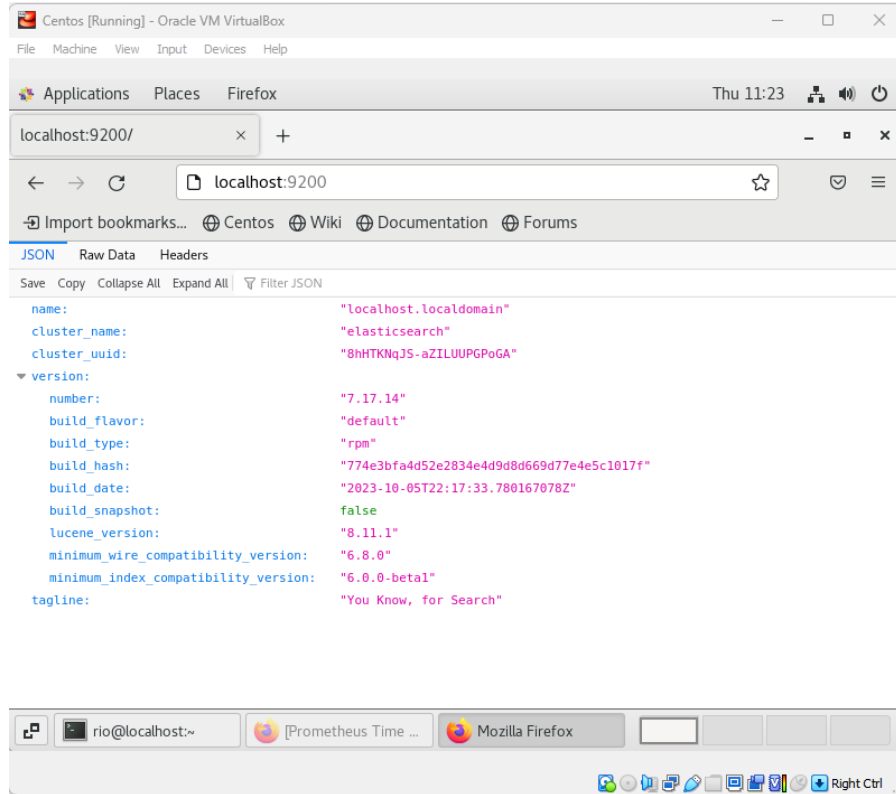
Output:

```
[rio@localhost ~]$ systemctl status logstash
● logstash.service - logstash
   Loaded: loaded (/etc/systemd/system/logstash.service; enabled; vendor preset: disabled)
   Active: active (running) since Thu 2023-10-12 11:20:10 EDT; 8s ago
     Main PID: 28103 (java)
       Tasks: 15
      CGroup: /system.slice/logstash.service
              └─28103 /usr/share/logstash/jdk/bin/java -Xms1g -Xmx1g -XX:+UseConcMarkSv...

Oct 12 11:20:10 localhost.localdomain systemd[1]: Started logstash.
Oct 12 11:20:10 localhost.localdomain logstash[28103]: Using bundled JDK: /usr/shar...k
Oct 12 11:20:10 localhost.localdomain logstash[28103]: OpenJDK 64-Bit Server VM war...
Hint: Some lines were ellipsized, use -l to show in full.
[rio@localhost ~]$
```







```
Server1 [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Activities Terminal Mon 19:10
rio@Server1: ~
File Edit View Search Terminal Help
valid_lft forever preferred_lft forever
rio@Server1:~$ systemctl status logstash
● logstash.service - logstash
   Loaded: loaded (/etc/systemd/system/logstash.service; enabled; vendor preset
   Active: active (running) since Mon 2023-10-23 19:08:42 PST; 10s ago
   Main PID: 6783 (java)
   Tasks: 15 (limit: 4656)
   CGroup: /system.slice/logstash.service
           └─6783 /usr/share/logstash/jdk/bin/java -Xms1g -Xmx1g -XX:+UseConcMa

Oct 23 19:08:42 Server1 systemd[1]: Started logstash.
Oct 23 19:08:42 Server1 logstash[6783]: Using bundled JDK: /usr/share/logstash/
Oct 23 19:08:43 Server1 logstash[6783]: OpenJDK 64-Bit Server VM warning: Optio

[1]+  Stopped                  systemctl status logstash
rio@Server1:~$ systemctl status kibana
● kibana.service - Kibana
   Loaded: loaded (/etc/systemd/system/kibana.service; enabled; vendor preset:
   Active: active (running) since Mon 2023-10-23 18:47:39 PST; 22min ago
   Docs: https://www.elastic.co
   Main PID: 2440 (node)
   Tasks: 11 (limit: 4656)
   CGroup: /system.slice/kibana.service
           └─2440 /usr/share/kibana/bin/./node/bin/node /usr/share/kibana/bin/

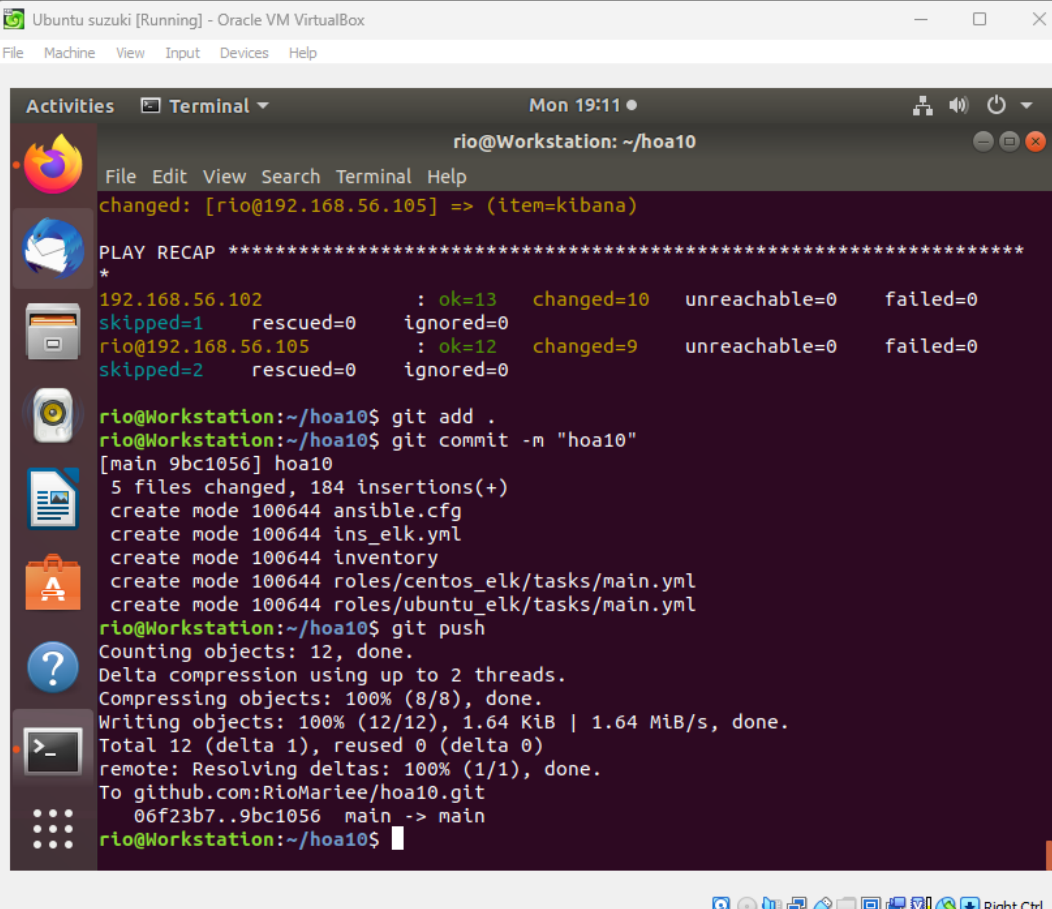
Oct 23 18:47:39 Server1 systemd[1]: Started Kibana.
Oct 23 18:47:39 Server1 kibana[2440]: Kibana is currently running with legacy 0
lines 1-11/11 (END)
```

```
Server1 [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Activities Terminal Mon 19:10
rio@Server1: ~
File Edit View Search Terminal Help
Loaded (/etc/systemd/system/kibana.service; enabled; vendor preset:
Active: active (running) since Mon 2023-10-23 18:47:39 PST; 22min ago
Docs: https://www.elastic.co
Main PID: 2440 (node)
Tasks: 11 (limit: 4656)
CGroup: /system.slice/kibana.service
        └─2440 /usr/share/kibana/bin/./node/bin/node /usr/share/kibana/bin/

Oct 23 18:47:39 Server1 systemd[1]: Started Kibana.
Oct 23 18:47:39 Server1 kibana[2440]: Kibana is currently running with legacy 0
lines 1-11/11 (END)
[2]+  Stopped                  systemctl status kibana
rio@Server1:~$ systemctl status elasticsearch
● elasticsearch.service - Elasticsearch
   Loaded: loaded (/usr/lib/systemd/system/elasticsearch.service; enabled; vend
   Active: active (running) since Mon 2023-10-23 18:47:37 PST; 22min ago
   Docs: https://www.elastic.co
   Main PID: 2095 (java)
   Tasks: 65 (limit: 4656)
   CGroup: /system.slice/elasticsearch.service
           └─2095 /usr/share/elasticsearch/jdk/bin/java -Xshare:auto -Des.netwo
           └─2302 /usr/share/elasticsearch/modules/x-pack-ml/platform/linux-x86

Oct 23 18:47:13 Server1 systemd[1]: Starting Elasticsearch...
Oct 23 18:47:18 Server1 systemd-entrypoint[2095]: Oct 23, 2023 6:47:18 PM sun.u
Oct 23 18:47:18 Server1 systemd-entrypoint[2095]: WARNING: COMPAT locale provid
Oct 23 18:47:37 Server1 systemd[1]: Started Elasticsearch.
lines 1-14/14 (END)
```

Step 11. After processing the `ins_elk` we will now commit our repository so that the code will be uploaded and updated in the github. The command we will use is “`git add .`” “`git commit -m "hoa10 "`” and “`git push`” to upload the latest version of the repository.



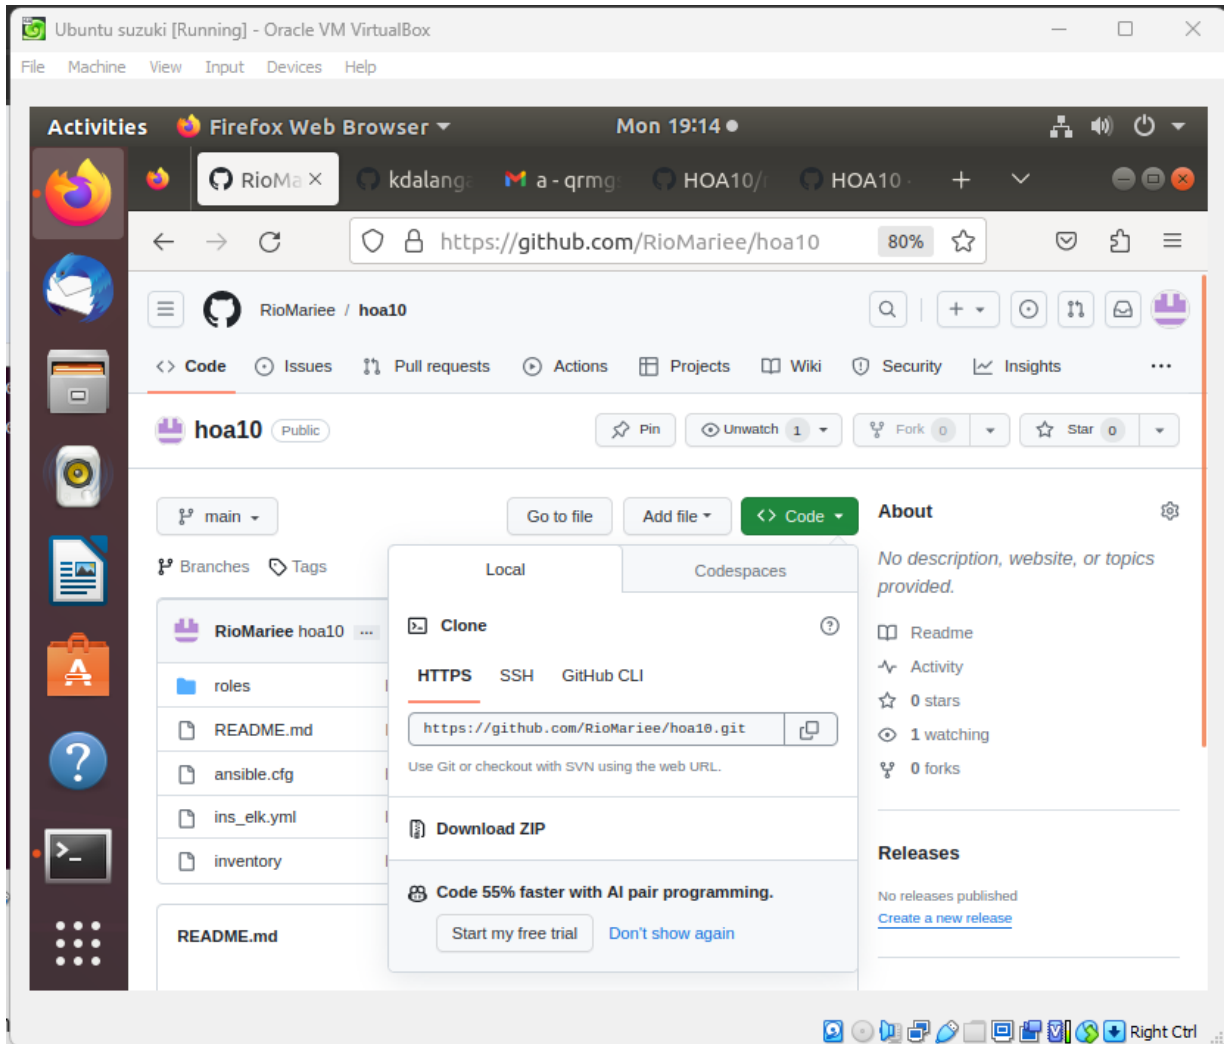
The screenshot shows a terminal window titled "Ubuntu suzuki [Running] - Oracle VM VirtualBox". The terminal output is as follows:

```
rio@Workstation: ~/hoa10
changed: [rio@192.168.56.105] => (item=kibana)

PLAY RECAP *****
*
192.168.56.102      : ok=13   changed=10   unreachable=0   failed=0
skipped=1         rescued=0   ignored=0
rio@192.168.56.105 : ok=12   changed=9    unreachable=0   failed=0
skipped=2         rescued=0   ignored=0

rio@Workstation:~/hoa10$ git add .
rio@Workstation:~/hoa10$ git commit -m "hoa10"
[main 9bc1056] hoa10
5 files changed, 184 insertions(+)
create mode 100644 ansible.cfg
create mode 100644 ins_elk.yml
create mode 100644 inventory
create mode 100644 roles/centos_elk/tasks/main.yml
create mode 100644 roles/ubuntu_elk/tasks/main.yml
rio@Workstation:~/hoa10$ git push
Counting objects: 12, done.
Delta compression using up to 2 threads.
Compressing objects: 100% (8/8), done.
Writing objects: 100% (12/12), 1.64 KiB | 1.64 MiB/s, done.
Total 12 (delta 1), reused 0 (delta 0)
remote: Resolving deltas: 100% (1/1), done.
To github.com:RioMarlee/hoa10.git
06f23b7..9bc1056  main -> main
rio@Workstation:~/hoa10$
```

Step 12. We will get our repository link and paste it here. To get the repository link go to github.com and click the repository you want to have the link in this case we will choose the hoa10 repository. Copy the http link and paste it here in the document.



Github repository link: <https://github.com/RioMarieee/hoa10.git>

Reflections:

Answer the following:

1. What are the benefits of having a log monitoring tool?
 - Log monitoring tools bring a lot of benefits for organizations in keeping their IT systems secure. These tools secure by spotting and tackling potential cyber threats quickly. They also help in making sure the systems run smoothly and don't crash. In addition, they help keep everything legal by keeping good

Conclusions:

Assessment Rubrics

Rubric for SO 7										Pts		
Criteria	Ratings								Pts			
SO 7 PI 1 Acquire and apply new knowledge from outside sources threshold: 4.8 pts	6 pts Excellent Educational interests and pursuits exist and flourish outside classroom requirements, knowledge and/or experiences are pursued independently and applies knowledge learned into practice		5 pts Good Educational interests and pursuits exist and flourish outside classroom requirements, knowledge and/or experiences are pursued independently		4 pts Satisfactory Look beyond classroom requirements, showing interest in pursuing knowledge independently		3 pts Unsatisfactory Begins to look beyond classroom requirements, showing interest in pursuing knowledge independently		2 pts Poor Relies on classroom instruction only	1 pts Very Poor No initiative or interest in acquiring new knowledge	6 pts	
SO 7 PI 2 Learn independently. threshold: 4.8 pts	6 pts Excellent Completes an assigned task independently and practices continuous improvement		5 pts Good Completes an assigned task without supervision or guidance		4 pts Satisfactory Requires minimal guidance to complete an assigned task		3 pts Unsatisfactory Requires detailed or step-by-step instructions to complete a task		2 pts Poor Shows little interest to complete a task independently		1 pts Very Poor No interest to complete a task independently	6 pts
SO 7 PI 3 Critical thinking in the broadest context of technological change threshold: 4.8 pts	6 pts Excellent Synthesizes and integrates information from a variety of sources; formulates a clear and precise perspective; draws appropriate conclusions		5 pts Good Evaluate information from a variety of sources; formulates a clear and precise perspective.		4 pts Satisfactory Analyze information from a variety of sources; formulates a clear and precise perspective.		3 pts Unsatisfactory Apply the gathered information to formulate the problem		2 pts Poor Gather and summarized the information from a variety of sources but failed to formulate the problem		1 pts Very Poor Gather information from a variety of sources	6 pts
SO 7 PI 4 Creativity and adaptability to new and emerging technologies threshold: 4.8 pts	6 pts Excellent Ideas are combined in original and creative ways in line with the new and emerging technology trends to solve a problem or address an issue.		5 pts Good Ideas are creative and adapt the new knowledge to solve a problem or address an issue		4 pts Satisfactory Ideas are creative in solving a problem, or address an issue		3 pts Unsatisfactory Shows some creative ways to solve the problem		2 pts Poor Shows initiative and attempt to develop creative ideas to solve the problem		1 pts Very Poor Ideas are copied or restated from the sources consulted	6 pts
Total Points: 24												